



**kafka**





## Module 2: Kafka Producer



# Objectives

---

After completing of this module, you should be able to:

- ✓ Configure Producers
- ✓ Construct Kafka Producer
- ✓ Send Messages to Kafka
- ✓ Synchronous and Asynchronous messages
- ✓ Serialize Messages using Avro
- ✓ Create and handle Partitions



# **Let's see how to configure Single-Node Multi-broker Kafka Cluster**



# Objectives

---

After completing of this module, you should be able to:

- ✓ Configure Producers
- ✓ Construct Kafka Producer
- ✓ Send Messages to Kafka
- ✓ Synchronous and Asynchronous messages
- ✓ Serialize Messages using Avro
- ✓ Create and handle Partitions



# Multi-Broker Cluster Setup

---

For setting up multiple brokers on a single node, separate server property files are required for each broker.  
Each property file will define different values for the following properties: *broker.id*, *listeners*, *log.dir*

## Steps:

- Go to Kafka directory
- Open *config* folder
- Make two separate server property files in config folder
- Make changes in the files created as shown on next slide



# Change in Server Property Files

server-1.properties

```
##### Server Basics #####  
# The id of the broker. This must be set to a unique integer for each broker.  
broker.id=1  
  
# Switch to enable topic deletion or not, default value is false  
#delete.topic.enable=true  
  
##### Socket Server Settings #####  
  
# The address the socket server listens on. It will get the value returned from  
# java.net.InetAddress.getCanonicalHostName() if not configured.  
# -DRVAL:  
# listeners = listener_name://host_name:port  
# EXAMPLE:  
# listeners = PLAINTEXT://your.host.name:9092  
listeners=PLAINTEXT://localhost:9093  
  
##### Log Basics #####  
  
# A comma separated list of directories under which to store log files  
log.dirs=/tmp/kafka-logs-1  
  
# The default number of log partitions per topic. More partitions allow greater  
# parallelism for consumption, but this will also result in more files across  
# the brokers.  
num.partitions=1
```

server-2.properties

```
##### Server Basics #####  
# The id of the broker. This must be set to a unique integer for each broker.  
broker.id=2  
  
# Switch to enable topic deletion or not, default value is false  
#delete.topic.enable=true  
  
##### Socket Server Settings #####  
  
# The address the socket server listens on. It will get the value returned from  
# java.net.InetAddress.getCanonicalHostName() if not configured.  
# -DRVAL:  
# listeners = listener_name://host_name:port  
# EXAMPLE:  
# listeners = PLAINTEXT://your.host.name:9092  
listeners=PLAINTEXT://localhost:9094  
  
##### Log Basics #####  
  
# A comma separated list of directories under which to store log files  
log.dirs=/tmp/kafka-logs-2  
  
# The default number of log partitions per topic. More partitions allow greater  
# parallelism for consumption, but this will also result in more files across  
# the brokers.  
num.partitions=1
```

Now we start each new broker in a separate console window



**Let's take a look at different Producer  
Application Use-cases**





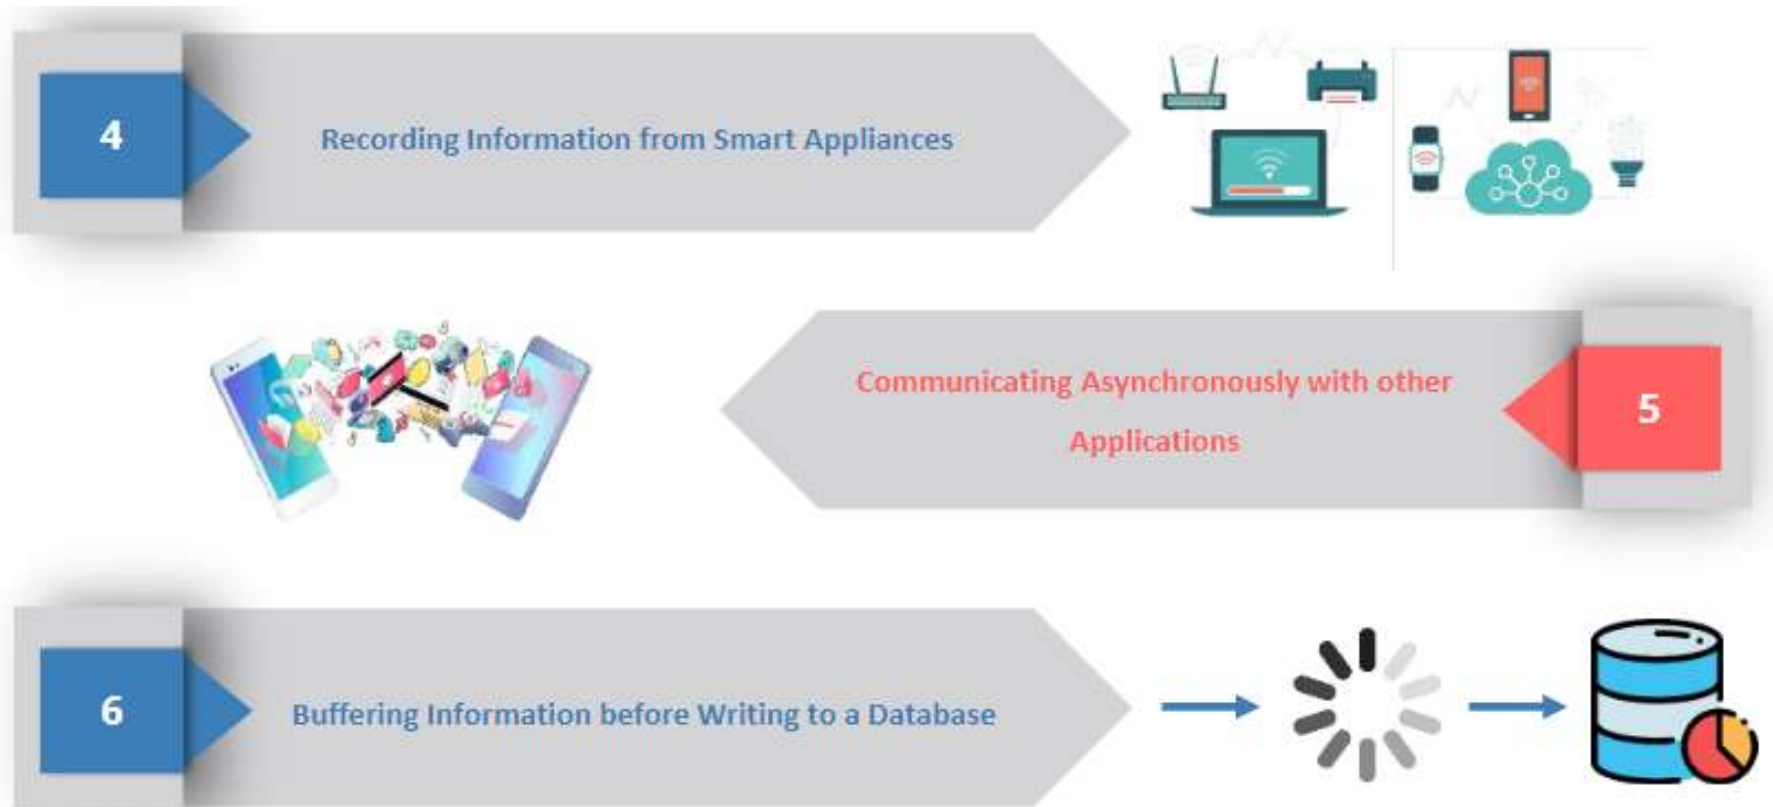
# Kafka Producer Applications Use-cases

---



# Kafka Producer Applications Use-cases

---



# **Let's take a look at different Kafka Producer Scenarios**



# Kafka Producer Scenarios

## Scenario 1: Credit Card Transaction Processing

- It is critical, since you never lose a single message nor duplicate any messages
- Latency should be low (can be tolerated up to 500ms)
- Throughput should be very high (process a million messages a second)



## Scenario 2: Clickstream Analysis

- Loss or a few duplicate messages can be tolerated
- Latency can be high as long as there is no impact on the user experience
- Throughput will depend on the level of activity



Different requirements will influence the way you use the producer API to write messages to Kafka and the configuration you use.



# High Level Architecture of Kafka Producer



# High Level Architecture of Kafka Producer

---

We can start producing messages to Kafka by creating  
a *ProducerRecord*

ProducerRecord



# High Level Architecture of Kafka Producer

---

It must include the *topic* we want to send the record to and a *value*



# High Level Architecture of Kafka Producer

---

We can also specify a *key* and/or a *partition*






# High Level Architecture of Kafka Producer

---


First the producer will *serialize* the key and value objects to *ByteArrays*



# High Level Architecture of Kafka Producer

---

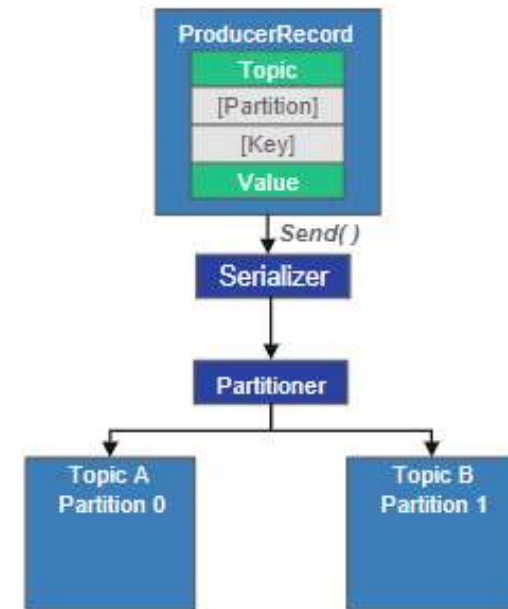
Next, the data is sent to a *partitioner*



# High Level Architecture of Kafka Producer

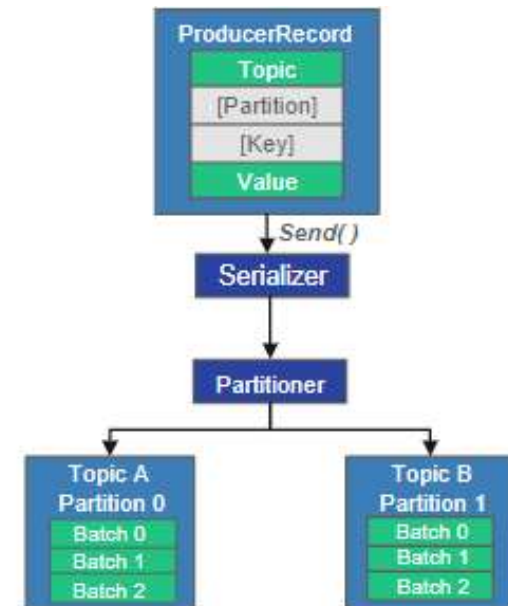
If partition is specified in *ProducerRecord*, the *partitioner* returns the *partition* we specified

As partition is selected, producer knows the topic and partition where the record will go

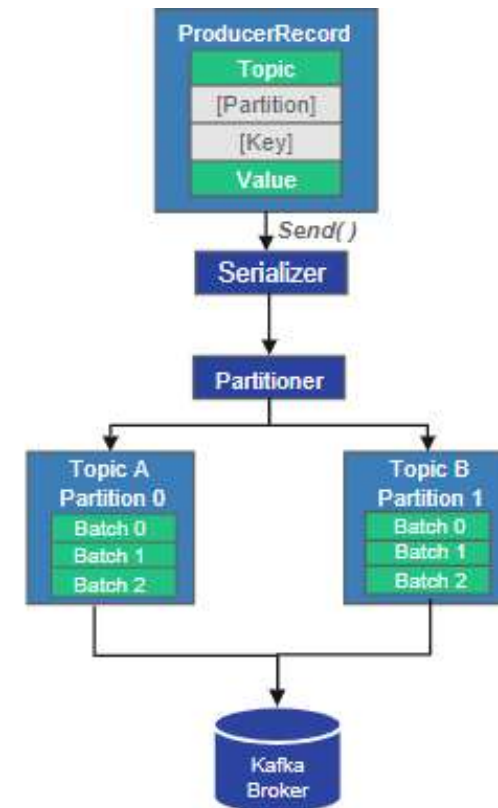


# High Level Architecture of Kafka Producer

Adds the record to a *batch of records* that will also be sent to the same topic and partition



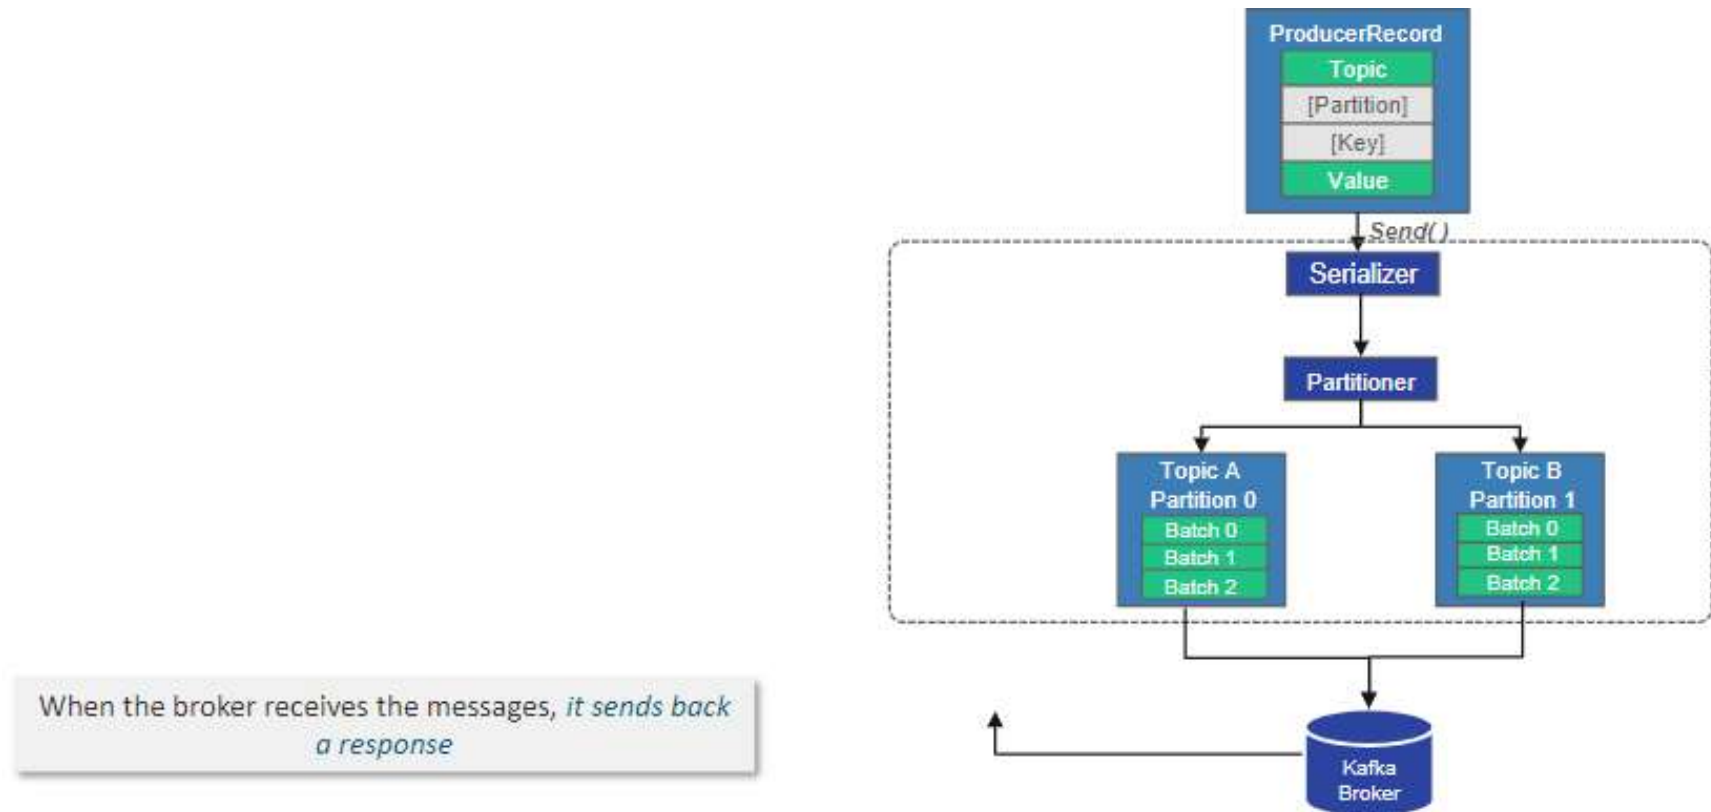
# High Level Architecture of Kafka Producer



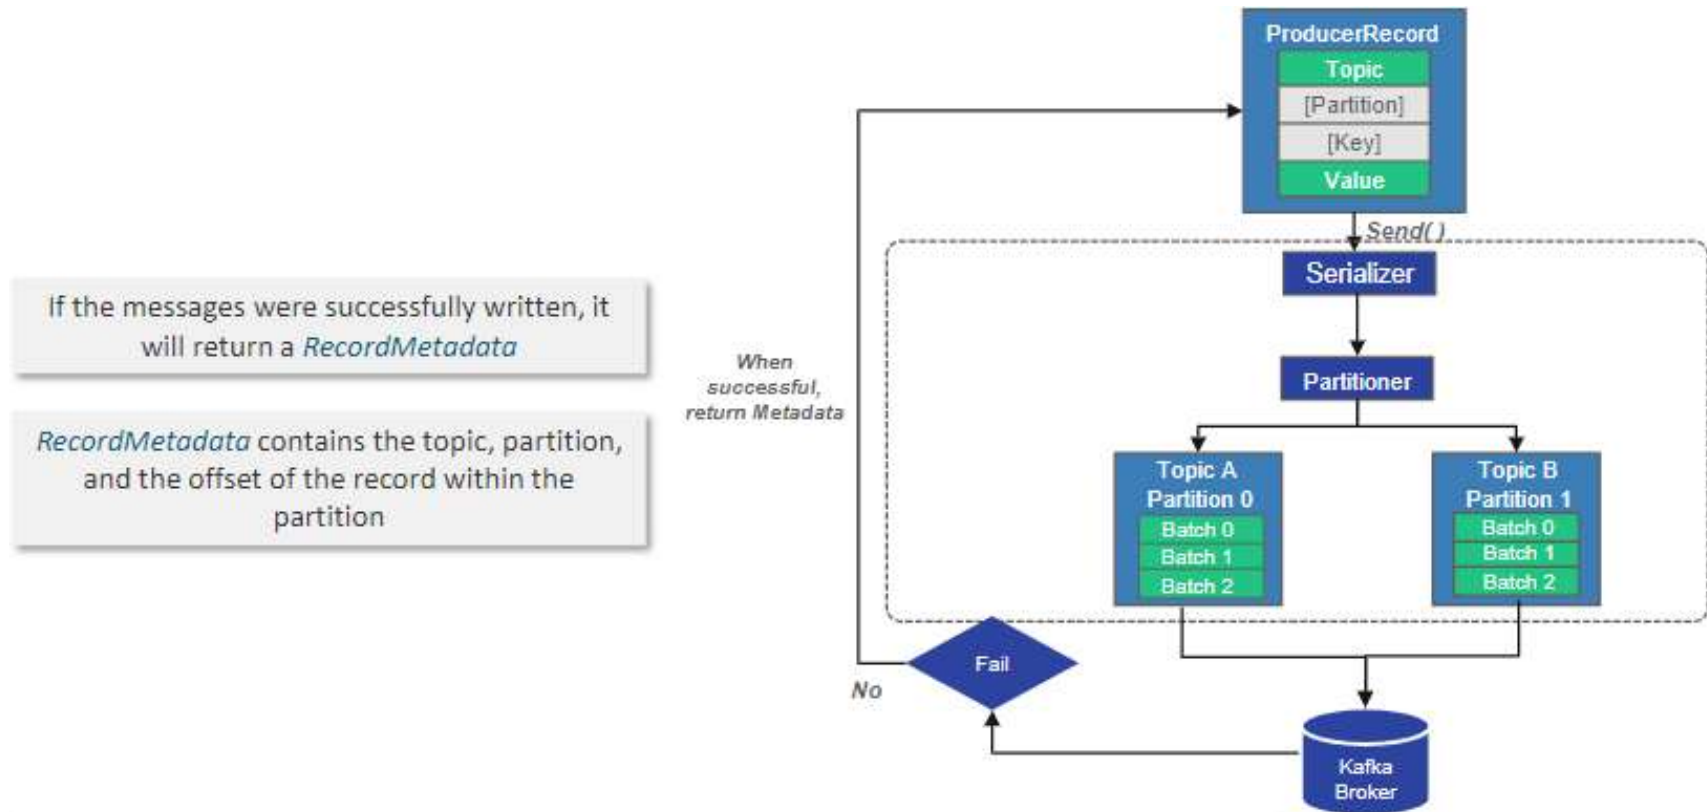
Separate thread is responsible for sending those batches of records to the appropriate *Kafka brokers*



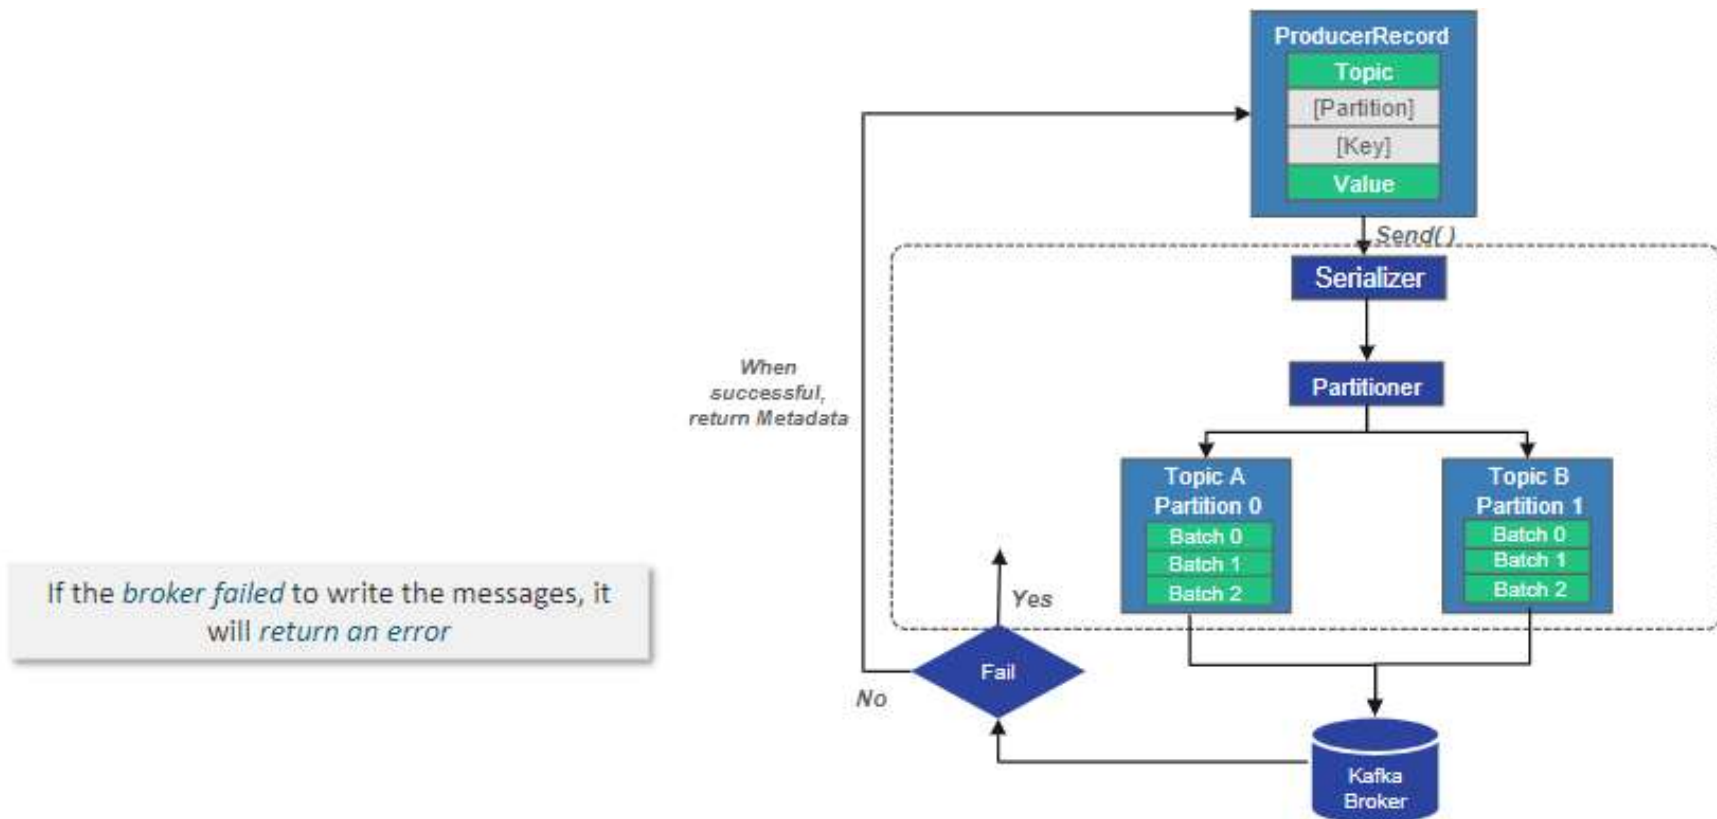
# High Level Architecture of Kafka Producer



# High Level Architecture of Kafka Producer

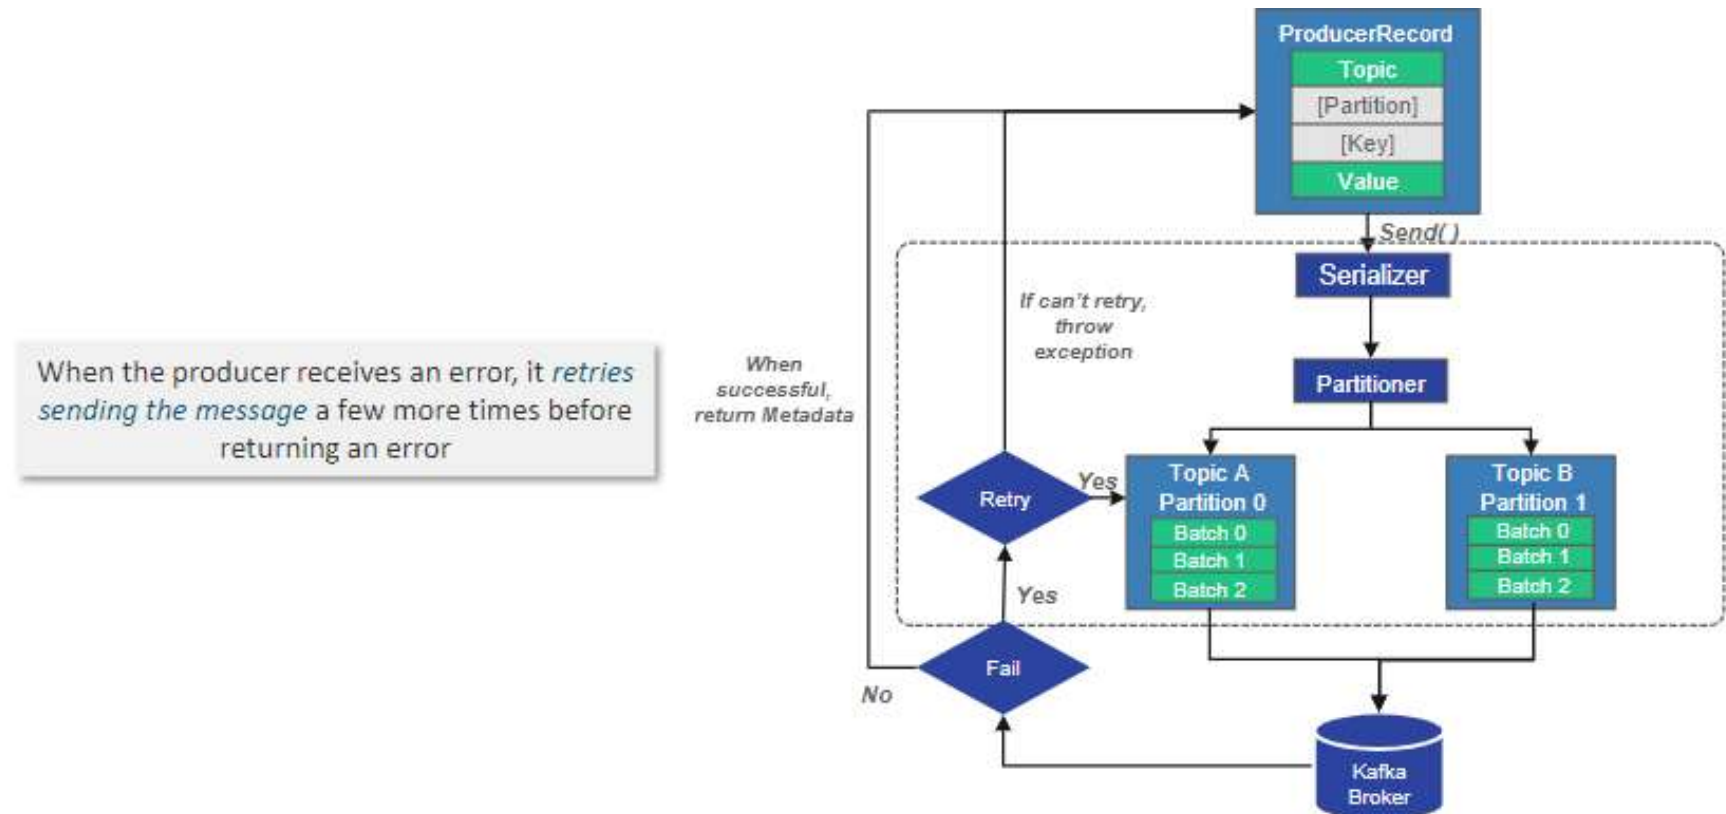


# High Level Architecture of Kafka Producer





# High Level Architecture of Kafka Producer



**Let's have a look at some important  
Kafka Producer Configuration Properties**



# Kafka Producer Configurations

---



Kafka producer has three mandatory properties:

***bootstrap.servers***

***key.serializer***

***value.serializer***

- List of *host:port* pairs of brokers, that the producer will use to establish initial connection
- No need to include all brokers, producer will get information after the initial connection
- Recommended to include at least two, in case one broker goes down



# Kafka Producer Configurations

---



Kafka producer has three mandatory properties:

*bootstrap.servers*

*key.serializer*

*value.serializer*

- Name of a class that will be used to serialize the keys of the records
- Brokers expect byte arrays as keys and values of messages
- Serializer class should implements *org.apache.kafka.common.serialization.Serializer* interface.
- Producer will use this class to serialize the key object to a byte array
- Kafka client package includes *ByteArraySerializer*, *StringSerializer*, and *IntegerSerializer*,
- Setting *key.serializer* is required even if you intend to send only values



# Kafka Producer Configurations

---



Kafka producer has three mandatory properties:

***bootstrap.servers***

***key.serializer***

***value.serializer***

- Name of a class that will be used to serialize the values of the records
- Similarly as you set *key.serializer* to a name of a class
- Set *value.serializer* to a class that will serialize the message value object



**Let's see how to create a Kafka Producer**



# Create a Kafka Producer

---

```
private Properties kafkaProps = new Properties();  
kafkaProps.put("bootstrap.servers", "broker1:9092, broker2:9093");  
kafkaProps.put("key.serializer",  
"org.apache.kafka.common.serialization.StringSerializer");  
kafkaProps.put("value.serializer",  
"org.apache.kafka.common.serialization.StringSerializer");  
producer = new KafkaProducer<String, String>(kafkaProps);
```

We start with  
a Properties object

Here we Create a new producer by setting  
the appropriate key and value types &  
passing the Properties object

Here we are using the  
built-in *StringSerializer*



**Let's take a look at different types of errors**





# Types of Errors

---

*KafkaProducer* has two types of errors:

## Retriable Error



- *Retriable* errors are those that can be resolved by sending the message again
  - For example, a connection error can be resolved by reestablishing a connection, “no leader” error can be resolved when a new leader is elected for the partition
- *KafkaProducer* can be configured to retry those errors automatically

## Non-Retriable Error



- Some errors will not be resolved by retrying
- In those cases, *KafkaProducer* will not attempt a retry & will return the exception immediately
  - For example, “message size too large”



**Let's see different ways to send messages**



# Ways to send Messages - Fire & Forget

---



Kafka producer has three ways for sending messages:

***Fire-and-forget***

***Synchronous send***

***Asynchronous send***

- We send a message to the server and don't really care if it arrives successfully or not
- Generally, messages arrive successfully, as Kafka is highly available
- Producer will retry sending messages automatically
- Some messages will get lost using this method



# Ways to send Messages - Synchronous Send

---



Kafka producer has three ways for sending messages:

*Fire-and-forget*

*Synchronous send*

*Asynchronous send*

- We send a message, the `send()` method returns a Future object, and we use `get()` to wait on the future and see if the `send()` was successful or not



# Ways to send Messages - Asynchronous Send

---



Kafka producer has three ways for sending messages:

***Fire-and-forget***

***Synchronous send***

***Asynchronous send***

- We call the `send()` method with a callback function, which gets triggered when it receives a response from the Kafka broker



# Sending Message in Fire & Forget way

---

```
ProducerRecord<String, String> record =  
    new ProducerRecord<>("Employee", "Name", "John");  
  
try {  
    producer.send(record); }  
  
catch (Exception e) {  
    e.printStackTrace(); }
```

- Producer accepts *ProducerRecord* objects
- *ProducerRecord* has multiple constructors
- Requires the name of the topic where we are sending data
- Always a string, key and value are also strings
- Key and value must match our serializer and producer objects



# Sending Message in Fire & Forget way

---

```
ProducerRecord<String, String> record =  
    new ProducerRecord<>("Employee", "Name", "John");  
  
try {  
    producer.send(record); }  
  
catch (Exception e) {  
    e.printStackTrace(); }
```

- Use producer object `send()` method to send the *ProducerRecord*
- Message will be placed in a buffer and will be sent to the broker in a separate thread
- *send()* method returns a Java Future object with *RecordMetadata*
- *RecordMetadata* tells whether the message was sent successfully or not



# Sending Message in Fire & Forget way

---

```
ProducerRecord<String, String> record =  
    new ProducerRecord<>("Employee", "Name", "John");  
  
try {  
    producer.send(record); }  
  
catch (Exception e) {  
    e.printStackTrace(); }
```

- Use producer object `send()` method to send the *ProducerRecord*
- Message will be placed in a buffer and will be sent to the broker in a separate thread
- *send()* method returns a Java Future object with *RecordMetadata*
- *RecordMetadata* tells whether the message was sent successfully or not





# Sending Message Synchronously

---

```
ProducerRecord<String, String> record = new ProducerRecord<>("Employee", "Name", "James");  
  
try {  
    producer.send(record).get();  
}  
  
catch (Exception e) {  
    e.printStackTrace();  
}
```

- *Future.get()* is used to wait for a reply from Kafka
- If the record is not sent successfully, method will throw an exception
- If there were no errors, it returns a *RecordMetadata* object



# Sending Message Synchronously

---

```
ProducerRecord<String, String> record = new ProducerRecord<>("Employee", "Name", "James");  
  
try {  
    producer.send(record).get();  
}  
  
catch (Exception e) {  
    e.printStackTrace();  
}
```

- It prints any exception, that has been occurred
- It can be errors before sending data, a nonretriable exceptions or available retries is exhausted



# Sending Message Asynchronously

```
private class DemoProducerCallback implements Callback {  
  
    @Override public void onCompletion(RecordMetadata recordMetadata, Exception e) {  
        if (e != null) {  
            e.printStackTrace();  
        }  
    }  
  
    ProducerRecord<String, String> record = new ProducerRecord<>("Employee", "Name", "Jordan");  
    producer.send(record, new DemoProducerCallback());  
}
```

- To use callbacks, a class is needed that implements the *org.apache.kafka.clients.producer.Callback* interface
- It has a single function—*onCompletion()*



# Sending Message Asynchronously

---

```
private class DemoProducerCallback implements Callback {  
  
    @Override public void onCompletion(RecordMetadata recordMetadata, Exception e) {  
        if (e != null) {  
            e.printStackTrace();  
        }  
    }  
  
    ProducerRecord<String, String> record = new ProducerRecord<>("Employee", "Name", "Jordan");  
    producer.send(record, new DemoProducerCallback());  
}
```

- If Kafka returned an error, *onCompletion()* will have a nonnull exception
- Production code will probably have more robust error handling functions



# Sending Message Asynchronously

```
private class DemoProducerCallback implements Callback {  
  
    @Override public void onCompletion(RecordMetadata recordMetadata, Exception e) {  
        if (e != null) {  
            e.printStackTrace();  
        }  
    }  
}  
  
ProducerRecord<String, String> record = new ProducerRecord<>("Employee", "Name", "Jordan");  
producer.send(record, new DemoProducerCallback());
```

- We pass a *Callback object* along when sending the record



**More properties to configure Kafka Producers ..**



# Configuring Kafka Producers

---

## ACKS

It controls how many partition replicas must receive the record before the producer can consider the write successful  
Significant impact on how likely messages are to be lost

*There are three allowed values for the acks parameter:*

### *acks=0*

- Producer will not wait for a reply from the broker before assuming the message was sent successfully
- If something went wrong and the broker did not receive the message
- Producer will not know about failure and the message will be lost
- Producer sends messages as fast as the network will support, it gives very high throughput

### *acks=1*

- Producer will receive a success response from the broker after leader replica receives the message
- If the message can't be written to the leader, the producer will receive an error response
- It can retry sending the message, avoiding potential loss of data
- Throughput depends on whether we send messages synchronously or asynchronously

### *acks=all*

- Producer will receive a success response from the broker once all in-sync replicas received the message
- It's the safest mode since you can make sure more than one broker has the message
- Message will survive even in the case of crash




# Configuring Kafka Producers

---

## buffer.memory

- Configures the amount of memory, producer will use to buffer messages before sending
- If messages are sent faster than they are delivered, producer may run out of space
- *Additional send() calls will either block or throw an exception (max.block.ms parameter)*



## compression.type

- By default, messages are sent uncompressed
- Compression algorithms will be used to compress the data before sending it to the brokers
- *It could be snappy, gzip, or lz4*





# Configuring Kafka Producers

---

## retries

- If error is transient, it tells how many times the producer will retry sending the message
- By default, the producer will wait 100ms between retries
- Could be controlled by *retry.backoff.ms* parameter



## client.id

- Used by the brokers to identify messages sent from the client
- Used in logging and metrics, and for quotas

## receive.buffer.bytes and send.buffer.bytes

- Sizes of the TCP send and receive buffers used by the sockets when writing/reading data
- If these are set to -1, the OS defaults will be used



# Configuring Kafka Producers

---

## `max.in.flight.requests.per.connection`

- Controls how many messages the producer will send to the server without receiving responses
- Setting this high can increase memory usage while improving throughput
- Setting it too high can reduce throughput as batching becomes less efficient
- Setting to 1 will guarantee that messages will be written to the broker in the order they were sent



## `request.timeout.ms`

- Controls how long the producer will wait for a reply from the server when sending data & requesting metadata
- If timeout is reached without reply, the producer will either retry sending or respond with an error



# Configuring Kafka Producers

## max.block.ms

- Controls how long the producer will block when calling `send()` & when explicitly requesting metadata via `partitionsFor()`
- Those methods block when the producer's send buffer is full or when metadata is not available
- When `max.block.ms` is reached, a timeout exception is thrown



## max.request.size

- This setting controls the size of a produce request
- Caps both the size of the largest message that can be sent & the size of messages in a batch that the producer can send in one request



# Serialization - What & How ?



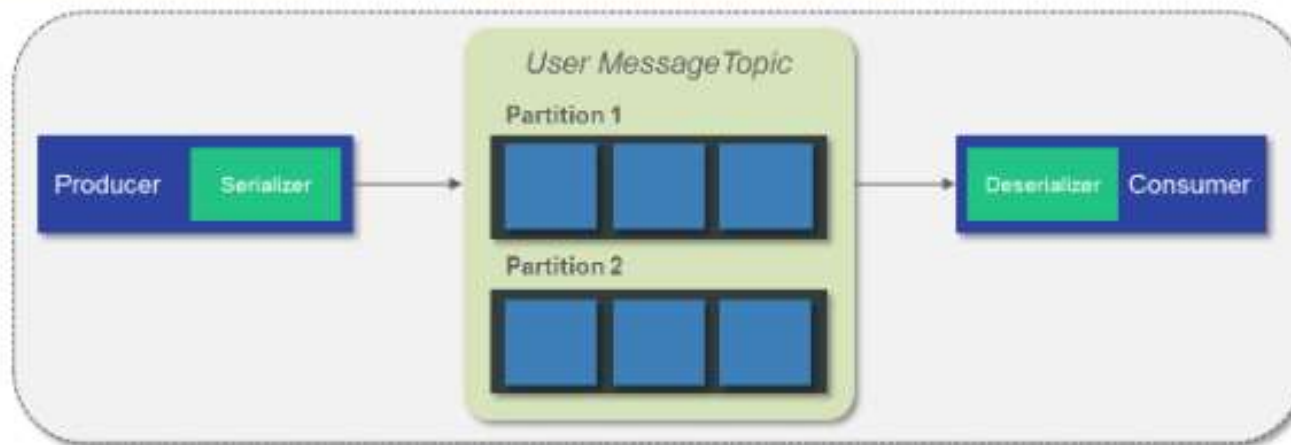
# Serializers

**Serialization** is the process of translating data structures or object state into a format that can be stored, transmitted & reconstructed later

Producer configuration includes mandatory serializers

Default String serializer can be used

Kafka also includes serializers for integers and ByteArrays



# Custom Serializers

---

## Custom Serializers

- If the object you need to send to Kafka is not a simple string or integer, you need *custom serializer*
- Can use generic serialization library like Avro, Thrift, or Protobuf to create records
- Can create a custom serialization for objects
- Creating a simple class to represent students:

```
public class Student {  
    private int studentID;  
    private String studentName;  
  
    public Student(int ID, String name) {  
        this.studentID = ID;  
        this.studentName = name; }  
  
    public int getID() {  
        return studentID; }  
  
    public String getName() {  
        return studentName; }  
}
```



# Creating Custom Serializer

---

E.g. - Creating a custom serializer for Student class

```
import org.apache.kafka.common.errors.SerializationException;
import java.nio.ByteBuffer;
import java.util.Map;

public class StudentSerializer implements Serializer<Student> {

    @Override
    public void configure(Map configs, boolean isKey) {
        // nothing to configure
    }

    @Override
    /**
    We are serializing Student as:
    4 byte int representing studentId
    4 byte int representing length of studentName in UTF-8 bytes (0 if name is Null)
    N bytes representing studentName in UTF-8
    */
}
```





# Creating Custom Serializer

---

E.g. - Creating a custom serializer for Student class

```
import org.apache.kafka.common.errors.SerializationException;
import java.nio.ByteBuffer;
import java.util.Map;
```

```
public class StudentSerializer implements Serializer<Student> {
```

```
    @Override
    public void configure(Map configs, boolean isKey) {
        // nothing to configure
    }
```

```
    @Override
    /**
```

```
    We are serializing Student as:
```

```
    4 byte int representing studentId
```

```
    4 byte int representing length of studentName in UTF-8 bytes (0 if name is Null)
```

```
    N bytes representing studentName in UTF-8
```

```
    */
```

- Configuring a producer with this *StudentSerializer* will allow you to define *ProducerRecord<String, Student>*
- Send *Student* data and pass *Student* objects directly to the producer





# Creating Custom Serializer

---

E.g. - Creating a custom serializer for Student class

```
public byte[] serialize(String topic, Student data) {
    try {
        byte[] serializedName;
        int stringSize;

        if (data == null)
            return null;
        else {
            if (data.getName() != null) {
                serializedName = data.getName().getBytes("UTF-8");
                stringSize = serializedName.length;
            } else {
                serializedName = new byte[0];
                stringSize = 0;
            }
        }
    }
}
```



# Creating Custom Serializer

---

E.g. - Creating a custom serializer for Student class

```
ByteBuffer buffer = ByteBuffer.allocate(4 + 4 + stringSize);
    buffer.putInt(data.getID());
    buffer.putInt(stringSize);
    buffer.put(serializedName);

    return buffer.array();
    } catch (Exception e) {
        throw new SerializationException("Error when serializing Student to byte[] " + e);
    }
}

@Override
public void close() {
    // nothing to close
}
}
```



# Serializer Challenges

If we need to change studentID to Long, or add a startDate field to *Student*, there will be compatibility issues between old and new messages.

Debugging compatibility issues is fairly challenging

If multiple teams are writing Student data to Kafka, they all need to use the same serializers & deserializer

It's recommended using existing serializers and deserializers such as JSON, Apache Avro, Thrift, or Protobuf



**Let's take a look at What is a Partition?**



# Partitions

---

- *ProducerRecord* objects includes a topic name, key & value
- Kafka messages are key-value pairs
- *ProduceRecord* object can be with just a topic and a value
- Key is set to null by default
- Most applications produce records with keys

Keys serve two goals:



Provides additional information that gets stored with the message

Used to decide which one of the topic partitions the message will be written to



# Partitions

---

- All messages with the same key will go to the same partition.
- All the records for a single key will be read by the same process.

To create a key-value record, you simply create a `ProducerRecord` as follows:

```
ProducerRecord<Integer, String> record = new ProducerRecord<>("Employee", "Name", "James");
```

While creating messages with a null key, you can simply leave the key out:

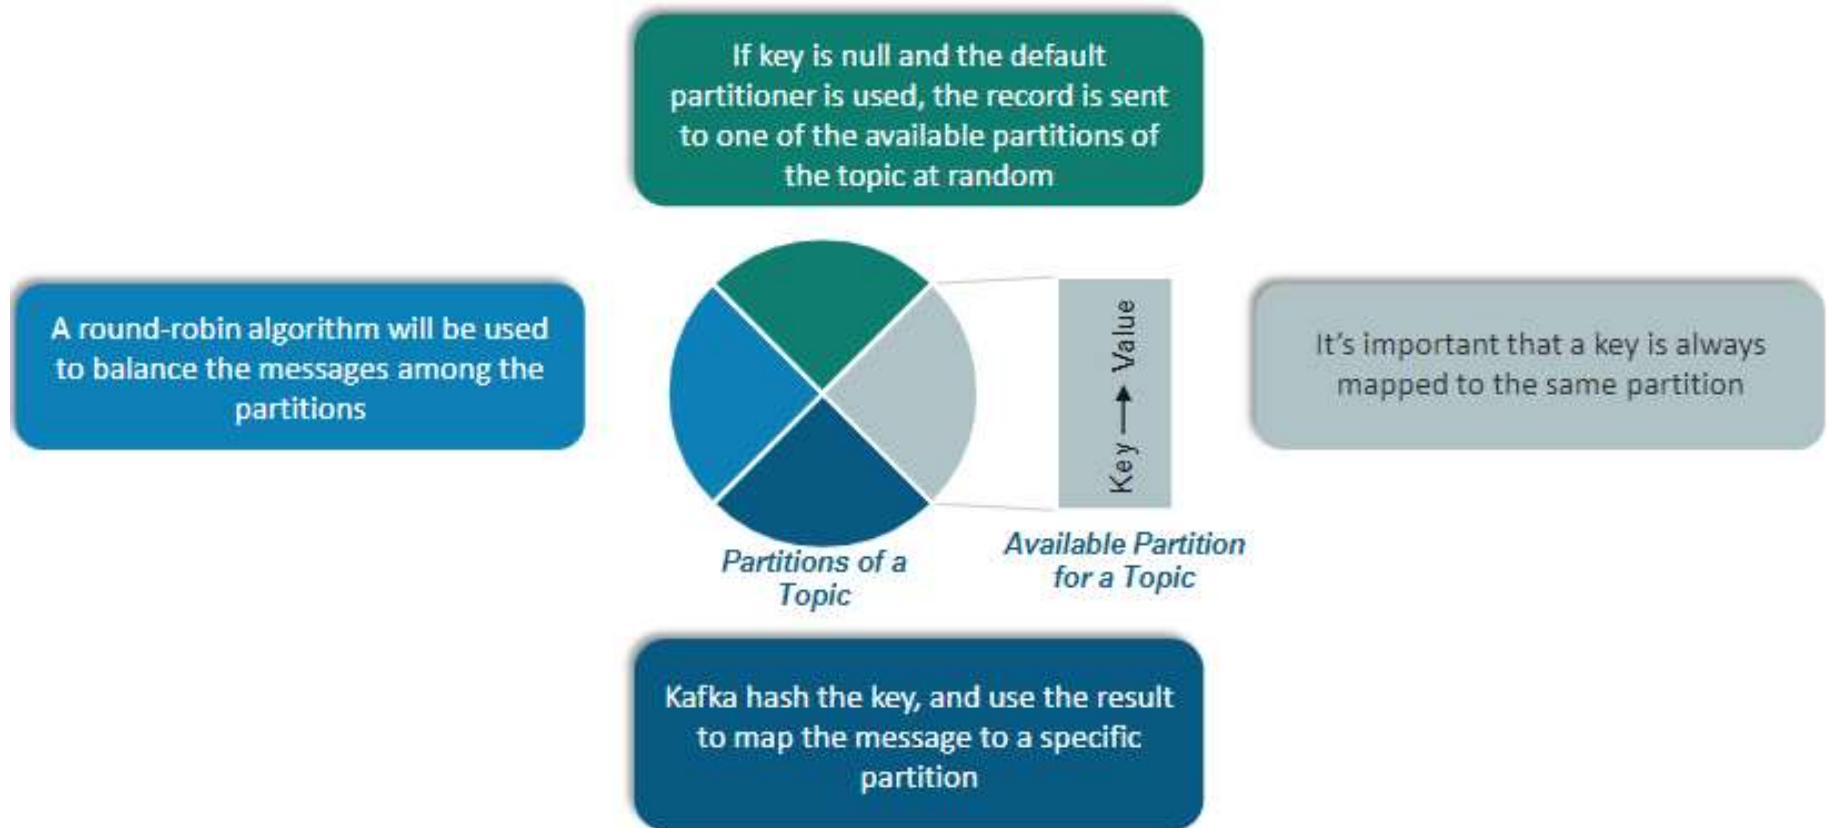
```
ProducerRecord<Integer, String> record = new ProducerRecord<>("Employee", "James");
```

Here, the key will simply be set to *null*, which may indicate that a student name was missing on a form.



# Partitions


---





# Partitions

---



As the number of partitions is constant, records of a user will be written in a particular partition

This allows all kinds of optimization when reading data from partitions



If you add *new partitions* to the topic, old records can change partition

New records will get written to a different partition



If partitioning keys is important, create topics with sufficient partitions and never add partitions





# Implementing a Custom Partitioning Strategy

Kafka allow to partition data differently.

## Scenario


- There is a company that does so much business with their device called "*Calacs*", that over 10% of their daily transactions are with this device.
- If you use default hash partitioning, the *Calacs* records will get allocated to the same partition as other accounts, resulting in one partition being about twice as large as the rest.

## Problem

This can cause servers to run out of space and slows down processing.

## Solution

To solve this problem, we need to provide *Calacs* its own partition and then use hash



*Calacs*



*Other devices*



*Sale*



# Implementing a Custom Partitioning Strategy

```
import org.apache.kafka.clients.producer.Partitioner;
import org.apache.kafka.common.Cluster;
import org.apache.kafka.common.PartitionInfo;
import org.apache.kafka.common.record.InvalidRecordException;
import org.apache.kafka.common.utils.Utils;
public class CalacsPartitioner implements Partitioner {
    public void configure(Map<String, ?> configs) {}
    public int partition(String topic, Object key, byte[] keyBytes, Object value, byte[]
valueBytes, Cluster cluster) {
        List<PartitionInfo> partitions = cluster.partitionsForTopic(topic);
        int numPartitions = partitions.size();
        if ((keyBytes == null) || (!(key instanceof String)))
            throw new InvalidRecordException("We expect all messages to have
            student name as key")
        if (((String) key).equals("Calacs"))
            return numPartitions; // Calacs will always go to last partition
        // Other records will get hashed to the rest of the partitions
        return (Math.abs(Utils.murmur2(keyBytes)) % (numPartitions - 1)) }

    public void close() {} }
}
```



# Implementing a Custom Partitioning Strategy

```
import org.apache.kafka.clients.producer.Partitioner;
import org.apache.kafka.common.Cluster;
import org.apache.kafka.common.PartitionInfo;
import org.apache.kafka.common.record.InvalidRecordException;
import org.apache.kafka.common.utils.Utls;
public class CalacsPartitioner implements Partitioner {
    public void configure(Map<String, ?> configs) {}
    public int partition(String topic, Object key, byte[] keyBytes, Object value, byte[]
valueBytes, Cluster cluster) {
        List<PartitionInfo> partitions = cluster.partitionsForTopic(topic);
        int numPartitions = partitions.size();
        if ((keyBytes == null) || (!(key instanceof String)))
            throw new InvalidRecordException("We expect all messages to have
            student name as key")
        if (((String) key).equals("Calacs"))
            return numPartitions; // Calacs will always go to last partition
        // Other records will get hashed to the rest of the partitions
        return (Math.abs(Utls.murmur2(keyBytes)) % (numPartitions - 1)) }

    public void close() {} }
```

Creating CalacsPartitioner



# Implementing a Custom Partitioning Strategy

```
import org.apache.kafka.clients.producer.Partitioner;
import org.apache.kafka.common.Cluster;
import org.apache.kafka.common.PartitionInfo;
import org.apache.kafka.common.record.InvalidRecordException;
import org.apache.kafka.common.utils.Utils;
public class CalacsPartitioner implements Partitioner {
    public void configure(Map<String, ?> configs) {}
    public int partition(String topic, Object key, byte[] keyBytes, Object value, byte[]
valueBytes, Cluster cluster) {
        List<PartitionInfo> partitions = cluster.partitionsForTopic(topic);
        int numPartitions = partitions.size();
        if ((keyBytes == null) || !(key instanceof String))
            throw new InvalidRecordException("We expect all messages to have
            student name as key")
        if (((String) key).equals("Calacs"))
            return numPartitions; // Calacs will always go to last partition
        // Other records will get hashed to the rest of the partitions
        return (Math.abs(Utils.murmur2(keyBytes)) % (numPartitions - 1)) }

    public void close() {} }
```

*Calacs will always go to last  
partition*





# Implementing a Custom Partitioning Strategy

```
import org.apache.kafka.clients.producer.Partitioner;
import org.apache.kafka.common.Cluster;
import org.apache.kafka.common.PartitionInfo;
import org.apache.kafka.common.record.InvalidRecordException;
import org.apache.kafka.common.utils.Utils;
public class CalacsPartitioner implements Partitioner {
    public void configure(Map<String, ?> configs) {}
    public int partition(String topic, Object key, byte[] keyBytes, Object value, byte[]
valueBytes, Cluster cluster) {
        List<PartitionInfo> partitions = cluster.partitionsForTopic(topic);
        int numPartitions = partitions.size();
        if ((keyBytes == null) || (!(key instanceof String)))
            throw new InvalidRecordException("We expect all messages to have
            student name as key")
        if (((String) key).equals("Calacs"))
            return numPartitions; // Calacs will always go to last partition
        // Other records will get hashed to the rest of the partitions
        return (Math.abs(Utils.murmur2(keyBytes)) % (numPartitions - 1)) }

    public void close() {} }
```

Other records will get hashed to the rest of the partitions



Thank you!

